Model building and validation

The formulas are primarily derived from *Linear Regression Analysis, Fifth Edition* by Montgomery et al. and/or the materials from the STAT 6021 course taught by Dr. Woo at the University of Virginia. Except where specially cited, it is believed that the materials are well-known equations and concepts in the public domain. If you believe otherwise, please reach out to me through my Github account so that I can correct the material.

Specifically, see chapters 10 and 11 of the textbook.

The best model

To quote the book,

> Unfortunately, as we will see in this chapter, there is no unique definition of "best." Furthermore, there are several algorithms that can be used for variable selection, and these procedures frequently specify different subsets of the candidate regressors as best.

Translation: You will need to read the textbook or other texts to study the specific issues and considerations that should inform your judgements.

Model adequacy versus validity

- Model adequacy concerns with whether the model and data meet the theoretical requirements, e.g., linearity assumption, constant variance assumption, etc.

- Model validation is about checking whether the model works well in practice, e.g., provides good predictions, works on new data not in the original estimation data set, etc.

Methods for variable selection

Comparing models

Coefficient of multiple determination

The value is

$$R_p^2 = \frac{SS_R(p)}{SS_T} = 1 - \frac{SS_Res(p)}{SS_T}$$

where $SS_R(p)$ is the regression sum of squares for the p-many coefficients ($p-1$ included regressors and the intercept), and SS_{Res} is the residual sum of squares for the same.

To find a "satisfactory" (to quote the book) value of R_p^2, you can follow one of the following approaches:

- Look for a convenient bend where the curve slopes off as p increases. This is a judgement call.
- Use the Atkins criteria (below).

Let

$$R_0^2 = 1 - (1 - R_{K+1}^2)(1 + d_{a,n,K})$$

where

$$d_{a,n,K} = \frac{KF_{a,K,n-K-1}}{n - K - 1}.$$

Then any value R_p^2 that exceeds R_0^2 is called an R^2-adequate (α) subset. Notice that the adequacy is based on a level of α. See p. 333 of the book for additional details and the reference to Atkins 1974.

Adjusted R^2

Let

$$R_{adj,p}^2 = 1 - \left(\frac{n-1}{n-p}\right)(1 - R_p^2)$$

be the adjusted version. You can look for when this stops increasing to see if models are ceasing to add sufficient additional value.

To quote the book:

> It can be shown (Edwards [1969]....) that if s regressors are added to the model, $R_{adj,p+s}^2$ will exceed $R_{adj,p}^2$ if and only if the partial F statistic for testing the significance of the s additional regressors exceeds 1.

Residual mean square

The residual mean square is

$$MS_{Res}(p) = \frac{SS_{Res}(p)}{n - p}$$

There are a few ways to pick the adequate model:

- The one with the minimum $MS_{Res}(p)$
- The one with an $MS_{Res}(p) \approx MS_{Res}^{full}$
- One near where changing adding/removing predictors causes an increase.

The book points out that where $MS_{Res}(p)$ is minimized, R_{adj}^2 will be maximized. See p. 334 for the proof.

Mallow's C_p statistic

Consider the mean square error of a single fitted value:

$$\text{MSE}(\hat{y}_i) = E[\hat{y}_i - E(y_i)]^2.$$

We can build on this (see p. 334) and define the standardized total mean squared error as

$$\Gamma_p = \frac{1}{\sigma^2}\left\{\sum_{i=1}^{n}[E(y_i) - E(\hat{y}_i)]^2 + \sum_{i=1}^{n}\text{Var}(\hat{y}_i)\right\}.$$

Assuming near-zero bias, we can estimate Γ_p as

$$\hat{\Gamma}_p = C_p = \frac{SS_{Res}(p)}{\hat{\sigma}^2} - n + 2p$$

where C_p is Mallows C_p statistic.

It can be shown that
$$E[C_p | \text{Bias} = 0] = p.$$

We can estimate $\hat{\sigma}^2 = MS_{Res}^{full}$ using the residual mean squared from the full model. This can be problematic if the full model has many unnecessary predictors. The text discusses an alternative.

Akaike Information Criterion (AIC)

The AIC is a "peanlized log-likelihood measure" that is derived from "maximizing the expected *entropy* of the model". (JG: Compare this to the concepts used in building decision trees based on entropy and/or the Gini coefficient).

It is defined (for least squares) as
$$AIC = n \ln(\frac{SS_{Res}}{n}) + 2p.$$

In the more general scenario, where L is the likelihood function of a particular model that isn't necessarily least squares, it is
$$AIC = -2 \ln(L) + 2p.$$

Bayesian information criteria (BICs)

There are two popular kinds of BICs.

The Schwartz BIC

It is defined as
$$\text{BIC}_{Sch} = -2 \ln(L) + p \ln(n),$$

which for least squares is
$$\text{BIC}_{Sch} = n \ln(\frac{SS_{Res}}{n}) + p \ln(n).$$

The Sawa BIC

We don't define it here.

Point estimators

The bias of an estimate is
$$B(\hat{\theta}) = E(\hat{\theta}) - \theta$$

The mean squared error for an estimate is
$$\text{MSE}(\hat{\theta}) = E((\hat{\theta} - \theta)^2).$$

It can be shown that
$$\text{MSE}(\hat{\theta}) = \text{Var}(\hat{\theta}) + (B(\hat{\theta}))^2.$$

See *Mathematical Statistics with Applications, ed. 6* by Wackerly et al. for additional details.

The text MSE is also used to represent the related concept of "mean square for error", discussed in the notebooks for single and multiple linear regression. Mean squared errors of point estimators is concerned with the theoretical values (assuming a known population), while mean square for error is normally calculated from samples.

Ancillary Material

The alias matrix

Let the full model be

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$$

Then in reducing the model to p many predictors by removing r many predictors, the $\mathbf{X}$ can be partitioned into two matrices, $\mathbf{X}_p$ and $\mathbf{X}_r$ such that the full model can be re-written as

$$\mathbf{y} = \mathbf{X}_p\boldsymbol{\beta}_p + \mathbf{X}_r\boldsymbol{\beta}_r + \boldsymbol{\varepsilon}.$$

Then

$$\mathbf{A} = (\mathbf{X}_p'\mathbf{X}_p)^{-1}\mathbf{X}_p'\mathbf{X}_r$$

is the alias matrix. The textbook describes its use, which we shall not repeat here.